
BUILDING A UNIX-TYPE FILE SYSTEM

OVERVIEW

This assignment focuses on UNIX-type file systems and their implementation. You will write a minimum file system using a simulated disk. Your file system will setup the data blocks to manage general information about the files, the i-nodes for file creation, and the free space.

THE PROGRAM

Create a program that sets up and formats a file system using a simulated disk as well as performs specific testing functions. The simulated disk reads and writes 1 KB binary data blocks from a random access file that is setup on your actual file system. Your file system consists of a super block, a list of i-node blocks, and data blocks. The super block is the first block (block 0) on disk. It contains information about the number of blocks on the file system, the number of i-nodes, and the index of the first block in the free list. The list of i-node blocks following the super block is needed to implement files. The following picture illustrates the layout of the disk:



DISK FORMATTING

Formatting a disk requires the following operations:

- 1) A super block is written to disk,
- 2) A list of empty i-nodes is created and written to disk, and
- 3) A free-space list is setup to indicate the data blocks that are free.

Review the code in *filesystem.c* and *filesystem.h* for functions that you must implement. These files contain the needed support to read and write integers from and to a byte array. In your implementation, you must choose the number of i-nodes that you want to store on the file system. Keep in mind that each file must have an i-node that describes the file's data blocks. This means that your file system must have as many i-nodes as you want files to be created.

FREE-SPACE MANAGEMENT

For free-space management, you will use a linked-list technique as discussed in class that links free blocks together. For example, if 100 is the first free block (as indicated by the superblock), the block contains the address (an integer) of the next free block, and so forth. The last free block contains -1 as the address of the free block to indicate the end of the free list.

TEST CODE

Write test code that prints formatting information to the screen after completion of the formatting process. Your test code must:

- 1) Display the information in the superblock, the number of i-nodes free and in use, and the number of disk blocks free and in use.
- 2) Perform consistency checking. This means that the test code must determine that the sum of the number of free disk blocks and the number of used disk blocks for data, superblock, and i-nodes is equal to the total number of disk blocks of the disk.
- 3) Create a binary file, write and read data of arbitrary size to the file, and close the file. The IO operations should closely resemble the equivalent IO operation of a regular file system. This means the following:
 - a) Files need to be created when they are opened for write operations and don't exist,
 - b) Open files grow in size with every write operation allowing the data to be appended to the existing data in the file,
 - c) Files cannot be opened for read operations if they don't exist,
 - d) Files exist in the directory after they are closed,
 - e) All open files are managed in a system-wide open file table.

Keep in mind that several independent write operations to a file must ensure that the file pointer moves along as bytes are written to the file indicating after every complete write operation the end of the file where the next byte can be written. And as data is written into the blocks, the last block of a file may only be partially filled.

IMPLEMENTATION SUGGESTIONS

To start this project, download the C files in the course shell in *eLearning*. Files *disk.c* and *disk.h* implement the simulated disk for reading and writing binary data blocks. Files *filesystem.c* and *filesystem.h* implement a rudimentary file system. The header files *superblock.h* and *inode.h* describe the superblock and i-node data structures. Both the superblock and a single i-node must each fit within a single disk block.

DELIVERABLES

Your project submission must follow the instructions below. Any submissions that do not follow the stated requirements will not be graded.

1. Follow the submission requirements of your instructor as published on *eLearning* under the Content area.
2. You must have at a minimum the following files for this assignment:
 - a. `disk.c`, `disk.h`
 - b. `fileSystem.c`, `fileSystem.h`
 - c. `superBlock.h`, `inode.h`, `inode.c`
 - d. `main.c`
 - e. `Makefile`
 - f. `README`

The README should describe any issues you may have encountered. If you do not include comments in your source code and refactor your code adequately, points will be deducted.

DUE DATE

The project is due as indicated by the Dropbox for project 4 in *eLearning*. Upload your complete solution to the dropbox and the shared drive. Upload ahead of time, as last minute uploads may fail. Please review the policy in the syllabus regarding late work.

TESTING

Test your code thoroughly on the public servers `ssh.cs.uwf.edu`. We will test your program on the testing server that uses the same OS and programming environment as the public servers.

GRADING

This project is worth 100 points total. The rubric used for grading is included below. Keep in mind that there will be deductions if your code does not compile, has memory leaks, or is otherwise, poorly documented or organized.

Submission	Perfect	Deficient		
eLearning	5 points individual files have been uploaded	0 points files are missing		
shared drive	5 points individual files have been uploaded	0 points files are missing		
Compilation	Perfect	Good	Attempted	Deficient
Makefile	5 points make file works; includes clean rule	3 points missing clean rule	2 points missing rules; doesn't compile project	0 points make file is missing
compilation	10 points no errors	7 points some warnings	3 points some errors	0 points many errors
Documentation & Style	Perfect	Good	Attempted	Deficient
documentation & program structure	5 points follows documentation and code structure guidelines	3 points follows mostly documentation and code structure guidelines; minor deviations	2 points some documentation and/or code structure lacks consistency	0 points missing or insufficient documentation and/or code structure is poor; review sample code and guidelines
File System Implementation	Perfect	Good	Attempted	Deficient
Creates i-nodes and superblock on disk	15 points correct, completed	11 points minor errors	4 points incomplete	0 points missing
Implements free space management	15 points correct, completed	11 points minor errors	4 points incomplete	0 points missing
Implements system call <code>openf()</code> , <code>readf()</code> , <code>writf()</code> , and <code>closef()</code>	20 points correct, completed	14 points minor errors	6 points incomplete	0 points missing
Implements test code	20 points correct, completed	14 points minor errors	6 points incomplete	0 points missing

I will evaluate your solution as attempted or insufficient if your code does not compile. This means, if you submit your solution according to my instructions, document and structure your code properly, and provide a makefile but the submitted code does not compile or crashes immediately you can expect at most 50 out of 100 points. So be sure your code compiles and executes properly.